


# Operating System

## Lab 2

### Matrix Multiplication (Multi-Threading)



Name : Mohamed Mohamed Hawas

ID : 18011595

# 1. Objectives

- To get familiar with thread programming using the Pthread library.
- To better understand processes and threads.

## 2. Overview

You are required to implement a multithreaded matrix multiplication program.

The input to the program is two matrixes  $A(xy)$  and  $B(yz)$  that are read from corresponding text files. The output is a matrix  $C(x*z)$  that is written to an output text file.

A parallelized version of matrix multiplication can be done using one of these three methods:

1. A thread computes the output C matrix i.e. without multi-threading. (A thread per matrix).

$$\begin{bmatrix} a_1 & a_2 & a_3 \\ a_4 & a_5 & a_6 \\ a_7 & a_8 & a_9 \end{bmatrix} \begin{bmatrix} b_1 & b_2 & b_3 \\ b_4 & b_5 & b_6 \\ b_7 & b_8 & b_9 \end{bmatrix} = \begin{bmatrix} c_1 & c_2 & c_3 \\ c_4 & c_5 & c_6 \\ c_7 & c_8 & c_9 \end{bmatrix}$$

2. A thread computes each row in the output C matrix. (A thread per row).

$$\begin{bmatrix} a_1 & a_2 & a_3 \\ a_4 & a_5 & a_6 \\ a_7 & a_8 & a_9 \end{bmatrix} \begin{bmatrix} b_1 & b_2 & b_3 \\ b_4 & b_5 & b_6 \\ b_7 & b_8 & b_9 \end{bmatrix} = \begin{bmatrix} c_1 & c_2 & c_3 \\ c_4 & c_5 & c_6 \\ c_7 & c_8 & c_9 \end{bmatrix}$$

3. A thread computes each element in the output C matrix. (A thread per element).

$$\begin{bmatrix} a_1 & a_2 & a_3 \\ a_4 & a_5 & a_6 \\ a_7 & a_8 & a_9 \end{bmatrix} \begin{bmatrix} b_1 & b_2 & b_3 \\ b_4 & b_5 & b_6 \\ b_7 & b_8 & b_9 \end{bmatrix} = \begin{bmatrix} c_1 & c_2 & c_3 \\ c_4 & c_5 & c_6 \\ c_7 & c_8 & c_9 \end{bmatrix}$$

### 3. Code Organization

The program reads two matrices from files, multiplies them using three different methods, and writes the resulting matrices to output files. The program uses “`pthread`s” to parallelize the multiplication process. Here is an overview of how the code is organized:

1. The program defines a `Matrix` struct that holds the dimensions and elements of a matrix.
2. It defines three instances of `Matrix`: `matrix_a`, `matrix_b`, and `matrix_c`.
3. The program defines three functions to read the contents of the matrices from files and populate the `Matrix` structs. These functions are `read_file`, `construct`, and `write_file`.
4. The program defines three multiplication functions:
  - `mult_per_matrix`: multiplies `matrix_a` and `matrix_b` using the matrix multiplication algorithm and stores the result in `matrix_c_per_matrix`.
  - `mult_per_row`: multiplies `matrix_a` and `matrix_b` using the matrix multiplication algorithm and stores the result in `matrix_c_per_row`. This function is designed to be run by multiple threads, with each thread computing one row of the result matrix.
  - `mult_per_element`: multiplies `matrix_a` and `matrix_b` using the matrix multiplication algorithm and stores the result in `matrix_c_per_element`. This function is designed to be run by

multiple threads, with each thread computing one element of the result matrix.

5. The main function reads the matrices from files using `read_file`, constructs the `matrix_c` matrices using `construct`, and creates the threads to perform the multiplications using `pthread_create`.
6. After the threads are created, the main function waits for all threads to complete using `pthread_join`.
7. Finally, the main function writes the resulting matrices to output files using `write_file`.

## 4. Code main functions

1. **`construct(Matrix* matrix, int rows, int cols)`**: A function that constructs a matrix with the given number of rows and columns.
2. **`void* read_file(void* arg, Matrix* matrix)`**: A thread function that reads the matrix data from a file and stores it in the given matrix.
3. **`void write_file(Matrix* matrix, char* filename)`**: A function that writes the matrix data to a file.
4. **`void* mult_per_matrix()`**: A thread function that multiplies two matrices (`matrix_a` and `matrix_b`) and stores the result in `matrix_c_per_matrix`.
5. **`void* mult_per_row(void* arg)`**: A thread function that multiplies two matrices (`matrix_a` and `matrix_b`) row by row and stores the result in `matrix_c_per_row`.
6. **`void* mult_per_element(void* arg)`**: A thread function that multiplies two matrices (`matrix_a` and `matrix_b`) element by element and stores the result in `matrix_c_per_element`.

### 3. Sample Runs and comparison

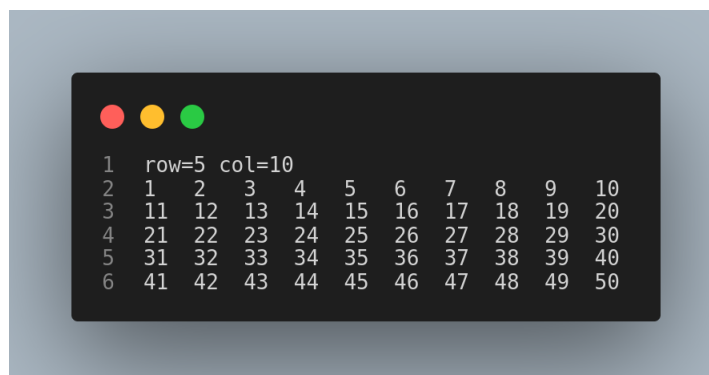
```
● hawas@hawas:~/Tech/OS/threads$ ./run.sh
Thread per matrix create 1 thread and take: 128 microseconds
Thread per row create 10 threads and take: 239 microseconds
Thread per element create 100 threads and take: 2281 microseconds
○ hawas@hawas:~/Tech/OS/threads$
```

- **a\_txt**



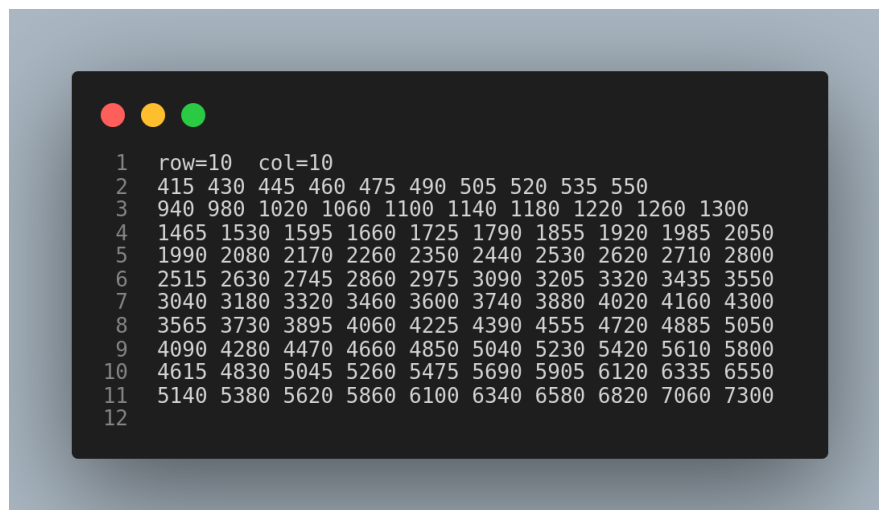
```
1 row=10 col=5
2 1 2 3 4 5
3 6 7 8 9 10
4 11 12 13 14 15
5 16 17 18 19 20
6 21 22 23 24 25
7 26 27 28 29 30
8 31 32 33 34 35
9 36 37 38 39 40
10 41 42 43 44 45
11 46 47 48 49 50
```

- **b\_txt**



```
1 row=5 col=10
2 1 2 3 4 5 6 7 8 9 10
3 11 12 13 14 15 16 17 18 19 20
4 21 22 23 24 25 26 27 28 29 30
5 31 32 33 34 35 36 37 38 39 40
6 41 42 43 44 45 46 47 48 49 50
```

- **c\_txt (For all)**



```
1 row=10 col=10
2 415 430 445 460 475 490 505 520 535 550
3 940 980 1020 1060 1100 1140 1180 1220 1260 1300
4 1465 1530 1595 1660 1725 1790 1855 1920 1985 2050
5 1990 2080 2170 2260 2350 2440 2530 2620 2710 2800
6 2515 2630 2745 2860 2975 3090 3205 3320 3435 3550
7 3040 3180 3320 3460 3600 3740 3880 4020 4160 4300
8 3565 3730 3895 4060 4225 4390 4555 4720 4885 5050
9 4090 4280 4470 4660 4850 5040 5230 5420 5610 5800
10 4615 4830 5045 5260 5475 5690 5905 6120 6335 6550
11 5140 5380 5620 5860 6100 6340 6580 6820 7060 7300
12
```